


# BUILD A DEEP CODE REVIEWER

## STEP-BY-STEP GUIDE



## BASHIRI SMITH

# Building a Deep Code Reviewer VS Code Extension – Step-by-Step Guide

This guide walks you through creating a “**Deep Code Reviewer**” – a Visual Studio Code extension powered by a Large Language Model (LLM) to review code for bugs, inconsistencies, and logical errors. We'll cover everything from setting up the development environment to integrating an LLM (like OpenAI's GPT-4 or Salesforce's CodeT5+), prompt design, fine-tuning the model, UI/UX considerations, testing, deployment, and real-world usage. The goal is a **professional-grade VS Code extension** that can be deployed for real-world development teams.

## 1. Environment Setup for VS Code Extension Development

Before coding the extension, set up your environment:

- **Install Node.js and Git:** VS Code extensions run in a Node.js environment. Ensure you have a recent Node.js (LTS) and Git installed ([Your First Extension | Visual Studio Code Extension API](#)).
- **VS Code and VS Code Extension SDK:** Download and install Visual Studio Code. You'll use VS Code's Extension API (the SDK is built-in with VS Code's `vscode` module).

**Yeoman Generator (Yo Code):** Microsoft provides a Yeoman generator to scaffold a new extension project. Install Yeoman and the VS Code Extension Generator globally with npm:

```
npm install --global yo generator-code
```

Alternatively, you can run it without global install using npx:

```
npx --package yo --package generator-code -- yo code
```

- Running `yo code` will prompt you for the extension's details and then create a basic project (with a Hello World command, a `package.json` manifest, and an entry point file) ([Your First Extension | Visual Studio Code Extension API](#)) ([Your First Extension | Visual Studio Code Extension API](#)). Choose “**New Extension (TypeScript)**” or “**New**

**Extension (JavaScript)**” when prompted for the type of extension.

- **Open and Run the Extension:** Open the generated project in VS Code. Press **F5** to launch a new Extension Development Host window (a sandbox VS Code instance for testing your extension). In that window, open the Command Palette (Ctrl+Shift+P) and run the “Hello World” command contributed by the sample. You should see a notification: “Hello World from <YourExtensionName>!” ([Your First Extension | Visual Studio Code Extension API](#)). This verifies the scaffold is working. If you don’t see the command, check that your `package.json`’s `engines.vscode` field matches your VS Code version ([Your First Extension | Visual Studio Code Extension API](#)).

**Tip:** The Yeoman scaffold also sets up a launch configuration for debugging. You can set breakpoints in your extension code (`extension.ts` or `extension.js`) and use VS Code’s debugger to step through your extension when you press F5 ([Your First Extension | Visual Studio Code Extension API](#)).

## 2. Choosing a Development Language: JavaScript vs. Python

VS Code extensions typically run in a Node.js context, so they are most often written in **TypeScript/JavaScript**. This offers seamless access to the VS Code Extension API and easy packaging via npm. If you are comfortable with TypeScript, it’s recommended for its type safety, but JavaScript is perfectly acceptable for extension development as well (the Yeoman tool can scaffold either).

What if you prefer Python? While you **cannot write a VS Code extension purely in Python**, you *can* use Python for the heavy lifting (such as machine learning tasks) by architecting your extension with a client-server model:

- **Node.js as the Extension Host:** The VS Code extension must be activated by VS Code’s Node.js-based runtime. This part can be minimal – just enough to capture user commands or editor events and then delegate processing.
- **Python for ML/LLM processing:** You can implement a separate Python service or utilize VS Code’s Language Server Protocol (LSP) to run a Python-based server. The extension (client) can spawn this Python process or communicate with it via IPC (inter-process communication). This approach is actually common: “*The Language Server can be written in any language*”, which is a key benefit of LSP ([Programmatic Language Features | Visual Studio Code Extension API](#)). For example, you might create a Python script that loads your fine-tuned model and listens on a port or stdin for code to review, then returns analysis results. Your extension’s JavaScript code can send the

code to this Python process and receive the feedback.

- **Pros and Cons:** Using JavaScript/TypeScript for everything keeps things simple within one process and is directly supported by VS Code's APIs. Using Python allows leveraging powerful ML libraries (TensorFlow/PyTorch) and existing code analysis tools, but introduces complexity in communication. If you opt for Python, plan for how the extension will manage the Python process (ensuring the correct environment, handling latency, etc.). Many real-world extensions (like Microsoft's Python and C++ extensions) use separate processes for heavy analysis, communicating via LSP or custom protocols.

**Summary:** For most developers, building the extension in JavaScript/TypeScript is easiest ([Your First Extension | Visual Studio Code Extension API](#)). You can still *incorporate* Python by having the JS extension call out to Python for the LLM evaluation. In this guide, we'll assume the core extension is in JavaScript/TypeScript, with notes on integrating Python if needed.

### 3. Selecting and Integrating a Large Language Model (LLM)

A "Deep Code Reviewer" relies on a capable LLM to analyze code. You have two broad choices:

- **Use a Hosted API LLM (e.g. OpenAI GPT-4):** This is the simplest route to get state-of-the-art performance. GPT-4 is a powerful model known for strong reasoning and code understanding. Through OpenAI's API, you send the code (and a prompt) and get back an analysis. This requires an API key and has usage cost, but no need to manage model infrastructure. OpenAI's GPT-4 is considered a leading model in 2025 for code tasks, especially if you have code review instructions in the prompt.
- **Use an Open-Source Code LLM (e.g. CodeT5+, Code Llama, StarCoder):** Open-source models can be run locally or on your own server. **CodeT5+** (by Salesforce) is a family of advanced code-focused LLMs that achieve state-of-the-art on many code tasks ([CodeT5+: Open Code Large Language Models](#)). Unlike general GPT models, CodeT5+ is designed to handle both *code understanding* (e.g. bug detection) and *code generation*, thanks to a flexible encoder-decoder architecture ([CodeT5+: Open Code Large Language Models](#)). Other notable models include Facebook's **Code Llama** (a variant of LLaMA tuned for coding) and **BigCode's StarCoder**. These models can often be run on-premises, which is beneficial for privacy (no sending code to an external API) and allows customization via fine-tuning.

**Embedding the LLM into the Extension:** There are a few integration patterns:



**Direct API Calls (Cloud LLM):** If using a service like OpenAI, integrate their REST API in your extension. For example, using the fetch API or an npm package, you would call the model with the code. **Code example (Node.js):**

```
const apiKey = process.env.OPENAI_API_KEY;
const prompt = `You are an expert software engineer... [include instructions & code here]`;
const response = await fetch("https://api.openai.com/v1/chat/completions", {
  method: "POST",
  headers: {
    "Content-Type": "application/json",
    "Authorization": `Bearer ${apiKey}`
  },
  body: JSON.stringify({
    model: "gpt-4",
    messages: [{role: "user", content: prompt}],
    temperature: 0 // deterministic output
  })
});
const result = await response.json();
const critique = result.choices[0].message.content;
```

- In this snippet, we send a chat prompt to GPT-4. The `messages` array format is used for OpenAI's chat API. We keep `temperature: 0` for consistent results. The `critique` will contain the model's analysis of the code (as a text response).
- **Local Model via REST or RPC:** If using an open model like CodeT5+ that you host, you might have a local server running (perhaps a FastAPI or Flask app in Python, or a Node HTTP server) that accepts code and returns the LLM's output. In the extension, you'd call `http://localhost:5000/review` (for example) with the code. Tools like **LM Studio** or the Hugging Face Inference API can serve models locally ([GitHub - hungson175/vscode-extension-gens: A gentle introduction to VS Extension & using LLM model](#)) ([GitHub - hungson175/vscode-extension-gens: A gentle introduction to VS Extension & using LLM model](#)). For instance, one tutorial uses LM Studio to host a 34B CodeLlama model and connects a VS Code extension to it via HTTP calls ([GitHub - hungson175/vscode-extension-gens: A gentle introduction to VS Extension & using LLM model](#)) ([GitHub - hungson175/vscode-extension-gens: A gentle introduction to VS Extension & using LLM model](#)). The pattern is similar: make a web request from the extension and receive the result.
- **Local Model via Python Process:** As discussed, you can also spawn a Python process from your extension. For example, use Node's `child_process.spawn` to run a Python script with the code as input, and capture the script's output (which contains the LLM's analysis). Ensure the Python environment has the model loaded (perhaps using Hugging Face Transformers to load CodeT5+ or Code Llama). This method avoids network calls

but requires managing the Python runtime.

**Working with Context Size:** Be mindful of the model's context window. GPT-4 and CodeT5+ can handle substantial input, but extremely large files might exceed limits. If needed, consider sending diffs or relevant sections rather than entire files. You can also chunk the code (e.g., analyze function by function) if the file is huge.

**Model Selection Tip:** For an MVP, using GPT-4 via API might get you quicker results (no ML ops overhead) and excellent quality. If you need offline usage or plan to fine-tune, using CodeT5+ or a similar open model is ideal. *In fact, CodeT5+ 16B has outperformed OpenAI's older code models on some benchmarks ([CodeT5+: Open Code Large Language Models](#)), indicating how capable open models have become.*

## 4. Prompt Engineering: Crafting Effective Code Review Prompts

Getting useful results from the LLM requires careful **prompt design**. You need to provide the model with the code and clear instructions on what to do. Here are some best practices for prompt engineering in this context:

- **Role and Task Description:** Start the prompt by setting the context. For example: *"You are a senior software engineer and code reviewer. Your task is to analyze the following code for any bugs, errors, or logical issues. Be thorough and explain any problems you find."* This gives the LLM a persona and a mission.

**Provide Code with Context:** Include the code to be reviewed in the prompt. If using OpenAI's chat format, you might do something like:

System: "You are an AI code reviewer. You identify bugs and issues in code and explain them."  
User: "<provide code here in a markdown block or fenced code block>"

- Make sure to delineate the code clearly (using Markdown formatting or a clear marker like `<CODE>` tags) so the model can distinguish instructions from code.
- **Specific Instructions for Output Format:** It's often helpful to ask the model to output in a structured way. This makes it easier for your extension to parse and display results. For instance, you can prompt: *"List each issue you find in the code. For each issue, provide a brief title, the affected line number(s), and an explanation. If possible, suggest a fix. If the code is correct, say there are no significant issues."* By specifying this, you increase the likelihood the model's response is organized (e.g., bullet points or a JSON structure). Using JSON is ideal for parsing – e.g., ask the model to output a JSON array of issues with fields `line`, `issue`, `severity`, `recommendation`. Some models follow

format instructions well, especially if you example it. (Note: Always include error-handling in case the model's output isn't perfectly formatted JSON).

- **Few-Shot Examples (if needed):** If the model has difficulty following instructions, you can include a **few-shot prompt** – i.e. give a small example of code and a sample review. For example: *“Example:\nCode:\npython\nx = 1\nif x = 2:\n print('equal')\n\nIssues:\n- Line 2: **Syntax Error** – Using a single equals sign in the *if* condition will cause a syntax error in Python. It should be *==* for comparison.”* Then follow with *“Now analyze the following code:”* and the real code. Few-shot prompting can guide the model on the style and depth you want.
- **Leverage Model-Specific Guidelines:** If using OpenAI, review their prompt engineering best practices ([Language Model API | Visual Studio Code Extension API](#)). For VS Code's built-in API (if you choose to use the new VS Code LanguageModel API), you can construct prompts using `LanguageModel1ChatMessage` and even define roles or personas for the model ([Language Model API | Visual Studio Code Extension API](#)).

**Testing your prompt:** Before wiring everything in the extension, it's useful to prototype your prompts in a standalone setting (for example, in the OpenAI Playground or using a curl command to the local model). Refine the prompt until the model's output is reasonably reliable in format and content.

## 5. Fine-Tuning the LLM to Catch More Subtle Mistakes (Optional, Advanced)

Out-of-the-box LLMs are generalists. Fine-tuning can teach the model about specific patterns of bugs or to follow a desired style of feedback. The question specifically mentions fine-tuning the model to catch mistakes *typical models might miss*. Fine-tuning is an advanced but powerful step:

- **Decide on a Model for Fine-Tuning:** OpenAI's GPT-4 is (at 2025) not openly fine-tunable by end-users, so for fine-tuning you'll likely use an open-source model like CodeT5+, Code Llama, or a smaller GPT-style model. For example, **LLaMA 3 8B** (an 8-billion parameter model) can be fine-tuned with techniques like LoRA to become a specialized code reviewer ([Fine-Tuning Small Language Models to Optimize Code Review Accuracy | NVIDIA Technical Blog](#)).
- **Gather a Training Dataset:** This dataset should consist of code samples *paired with* reviews (i.e., bug annotations or comments). You might source this from:
  - Real code reviews: If you have access to a repository of code reviews (e.g., GitHub pull request comments or Gerrit reviews), extract code diffs and review

comments as training pairs.

- Synthetic bugs: Intentionally introduce common mistakes into code and label them. For instance, create a function with an off-by-one error and write the expected critique pointing that out.
- Public datasets: Research and academia have some resources. For example, Google's DeepMind has a Code Review dataset, or you can use bug databases (ManyBugs, Defects4J for Java, etc.) where code with known bugs and fixes are available (the "fix" can guide what the bug was).
- **Fine-Tuning Process:** Use a framework like Hugging Face Transformers. Format your data as prompt→response. For example, the prompt could be the code (and possibly instructions), and the response is the ideal review (list of issues). Then perform supervised fine-tuning. Given the model size, you may use **LoRA (Low-Rank Adaptation)** to fine-tune efficiently on consumer hardware. NVIDIA's research shows that using LoRA on a smaller model can significantly boost its accuracy on code review tasks – in one case a fine-tuned 8B model outperformed a 70B model on bug detection after a ~18% accuracy improvement from fine-tuning ([Fine-Tuning Small Language Models to Optimize Code Review Accuracy | NVIDIA Technical Blog](#)) ([Fine-Tuning Small Language Models to Optimize Code Review Accuracy | NVIDIA Technical Blog](#)).
- **Evaluate the Fine-Tuned Model:** After training, test your model on some held-out code with known issues. If you have metrics like precision/recall of bug detection, measure them. Fine-tuning can dramatically improve targeted skills; OpenAI reported that a GPT-4-based "*critic model*" (fine-tuned specifically to critique code) was able to catch more bugs than human reviewers in some cases ( ) ( ). This indicates a well-tuned model can find subtle issues that a general model or even humans might miss. However, be cautious: fine-tuned models can sometimes "hallucinate" issues that aren't real ( ), so thorough testing is important.
- **Integrate the Fine-Tuned Model:** Deployment of your custom model could be via a local inference server. For example, if you fine-tuned a CodeT5+ model and saved it to Hugging Face Hub or to disk, you can load it in a Python service and query it. Ensure your extension can easily switch between using the general model (like GPT-4 API) and your fine-tuned model (maybe via a setting). For instance, "*modelProvider*": "*OpenAI*" vs "*LocalFineTuned*" in a config.

Fine-tuning is optional – you can get a lot of mileage from prompting alone – but it can push the performance further and tailor the LLM to consistently catch certain classes of bugs (for example, security vulnerabilities) that you emphasize in training. It's a **real-world deployment consideration** if you want the best accuracy and have resources to train models.



## 6. Implementing the Extension Logic – Commands, LLM Calls, and Results

With the environment ready, model chosen, and prompts in hand, now implement the VS Code extension's functionality. We'll outline a typical workflow: when the user invokes the “Deep Code Review” command, the extension will collect the code, send it to the LLM, then present the results in VS Code.

**6.1 Registering a Command:** In your extension's activation file (`extension.ts` or `.js`), use the VS Code API to register a command in the extension's activation function. For example:

```
import * as vscode from 'vscode';

export function activate(context: vscode.ExtensionContext) {
  const disposable = vscode.commands.registerCommand('deepCodeReviewer.reviewCode',
  async () => {
    // Command implementation will go here
  });
  context.subscriptions.push(disposable);
}
```

Also ensure you declared this command in your `package.json` (under the `"contributes"`: `{ "commands": [ ... ] }` section, with an identifier like `deepCodeReviewer.reviewCode` and a title).

You might also add a keyboard shortcut (in `contributes.keybindings`) and a context menu entry so that the command can be invoked easily (e.g., right-click in editor -> “Deep Code Review”).

**6.2 Gathering the Code Context:** Inside the command's callback, determine what code to review:

If you want to review the entire file, get the active text editor and retrieve its document text:

```
const editor = vscode.window.activeTextEditor;
if (!editor) {
  vscode.window.showInformationMessage("No active editor - open a file to review.");
  return;
}
const code = editor.document.getText();
const language = editor.document.languageId; // e.g. 'javascript', 'python'
```

-

- If you prefer to review only a selection (optional feature), you can use `editor.selection` to get the selected text (or default to whole document if nothing selected).
- You might also gather some metadata – like filename or language – and include that in the prompt (e.g., “The following is a JavaScript file named App.js:” to give model more context).

**6.3 Sending the Code to the LLM:** Use the integration method you chose in Step 3. For demonstration, assume we call an API (like OpenAI):

```
vscode.window.withProgress({ location: vscode.ProgressLocation.Notification, title: "Deep Code Review: Analyzing code...", cancellable: false }, async () => {
  try {
    const prompt = makeReviewPrompt(code, language); // function you write to insert code into your prompt template
    const critique = await callLLMApi(prompt);        // function that calls GPT-4 or your model
    handleLLMResponse(critique);
  } catch (err) {
    vscode.window.showErrorMessage("Code review failed: " + String(err));
    console.error(err);
  }
});
```

Here we wrap the call in `withProgress` to show a spinner notification in VS Code so the user knows analysis is happening (important because LLM calls might take a few seconds or more). We then handle errors gracefully, showing a message if something went wrong (network error, etc.).

**6.4 Processing the LLM's Response:** The `critique` we get back is presumably a structured text (or JSON). Now we need to present it to the user. There are a few ways to do this:

- **As a Text Output (Simple):** You can simply display the response in a VS Code Output Channel or a pop-up:
  - *Output Channel:* Create one if not exists: `const output = vscode.window.createOutputChannel("Code Review Results");` and then do `output.clear(); output.appendLine(critique); output.show();`. This will open a panel with the raw text. While easy, this is a bit basic and not interactive.

- *Information Message*: If the response is short, you could use `vscode.window.showInformationMessage`, but code reviews are often lengthy – not suitable for an alert box.

**As Editor Annotations (Preferred for issues)**: A polished approach is to mark up the code with the findings. VS Code’s Diagnostics API lets you create problem markers in the code (similar to compile errors or linter warnings). You can parse the LLM’s output for line numbers and messages, then create diagnostics so they show up in the “Problems” panel and as squiggly underlines in the code editor. For example, if the model output (structured) says there’s an issue at line 10, you can do:

```
const diagnostics: vscode.Diagnostic[] = [];
issues.forEach(issue => {
  const lineNum = issue.line - 1; // zero-based index
  const range = new vscode.Range(lineNum, 0, lineNum, Number.MAX_SAFE_INTEGER);
  const diagnostic = new vscode.Diagnostic(range, issue.message,
vscode.DiagnosticSeverity.Warning);
  diagnostics.push(diagnostic);
});
const collection = vscode.languages.createDiagnosticCollection('deep-code-review');
collection.set(editor.document.uri, diagnostics);
```

- This pseudo-code iterates over issues and creates a Diagnostic for each ([Programmatic Language Features | Visual Studio Code Extension API](#)). We use the whole line range for simplicity (or if we have column info, we can mark the specific substring). We add all diagnostics to a collection named `'deep-code-review'`. Once set, VS Code will underline the text and list the issues in the Problems panel. The user can click an issue in the Problems panel to jump to that line.

Using diagnostics integrates smoothly with VS Code’s UX for showing issues. The figure below shows an example of VS Code diagnostics in action (red underlines and messages for code issues):

([Programmatic Language Features | Visual Studio Code Extension API](#)) ([Programmatic Language Features | Visual Studio Code Extension API](#))

- **Code Lens or Hover (Optional Enhancements)**: For a more interactive feel, you can use CodeLens to show an annotation above the code (like “3 issues found by Deep Code Reviewer” with a link to details) ([Programmatic Language Features | Visual Studio Code Extension API](#)). Or implement a HoverProvider so that hovering over an underlined issue shows the LLM’s explanation in a tooltip. These use VS Code language features API (`registerCodeLensProvider`, `registerHoverProvider`) – not mandatory,

but nice to have if you want to mimic a real-time assistant feel.

- **Webview Panel (Rich UI):** If you want a custom UI (with formatting, maybe even a diff view for suggested fixes), you could create a Webview panel. For instance, display the code on one side and issues on the other, or have a styled list of issue cards. This requires more work (building an HTML/JS app within the panel), so consider it if you need a very specific UX. In many cases, using the editor's existing decoration mechanisms (diagnostics, hovers) is sufficient and consistent with how developers see errors from linters or compilers.

After deciding on the display approach, implement the `handleLLMResponse` accordingly. For example, if the LLM returned JSON, parse it and create diagnostics. If it returned markdown text, you might parse line numbers by regex. (Training the model to output JSON can simplify this step significantly.)

**6.5 Providing Fixes or Quick Actions (Optional):** As a bonus feature, if the LLM suggests fixes, you can use VS Code's Code Actions API to offer quick fixes. For example, if the model said "Line 10: off-by-one error, change `<=` to `<`", you can programmatically apply that edit via the `TextEditor` edit API when the user accepts the fix. This turns your tool into not just a reviewer but an auto-fixer (similar to ESLint's "fix" functionality). Implementing code actions would involve `registerCodeActionsProvider` and supplying `CodeAction` objects with edits ([Programmatic Language Features | Visual Studio Code Extension API](#)).

For the scope of this guide, we'll note that this is possible, but it requires reliable parse of suggested changes. You might start by just pointing out issues, and add fix support in a later iteration.

## 7. UI/UX Design Tips for a VS Code-integrated Tool

Designing a good UI/UX for a code review assistant is crucial for developer adoption. Here are some tips to make the integration seamless and developer-friendly:

- **Non-intrusive Operation:** Developers don't want constant distractions. It's wise to make the code review run **on-demand** (e.g., via command or button click) rather than automatically on every code change. This gives the user control, especially since LLM calls can be slow or costly. You might integrate it with VS Code's save or commit events only if performance is optimized. Start with an explicit trigger.
- **Leverage Familiar UI Elements:** As described above, using the Problems panel and editor decorations makes your tool feel like a natural extension of the editor. For example, a misused API flagged by the LLM will appear just like a compiler error in VS Code – highlighted in the code and listed in the Problems view. This consistency means

less learning curve for users.

- **Clarity and Detail Balance:** In the LLM's output, ensure the messages are clear and concise. An issue's description should ideally fit in one or two sentences, focusing on *what's wrong and why*. If the model's explanation is too verbose, consider trimming or summarizing it before displaying. However, do keep enough detail that the user learns from the feedback (the value of AI is often in the rationale it provides).
- **Severity Indication:** Not all issues are equal. You can map certain keywords or categories from the LLM's findings to severity levels. For instance, an actual bug or error would be **Error** severity, a code smell or style issue might be **Warning**, and a minor suggestion could be **Information**. Setting the `vscode.DiagnosticSeverity` appropriately will show different icons (e.g., x or exclamation) ([Programmatic Language Features | Visual Studio Code Extension API](#)). This helps developers prioritize.
- **Loading Feedback:** Always give the user a visual cue when analysis is running (using `withProgress` as shown, or a status bar item). If an operation might take 10+ seconds (say analyzing a very large file with a slower local model), inform the user of progress or an estimate. Maybe use a status bar message like "Deep Code Reviewer: analyzing..."
- **Configurability:** Provide a few useful settings in your extension's `package.json` configuration. For example, allow users to select which model or API to use (OpenAI vs local), an API key input (store securely via VS Code's Secrets API if possible), whether to run on save automatically, etc. Good defaults with the ability to change will make your extension useful to a wider audience.
- **Testing on Multiple Languages:** If your extension is meant to review code in multiple languages (Python, JavaScript, etc.), test the UI with each. Ensure that the Problems panel correctly scopes issues by file and that your prompts adapt (you might have the prompt mention the language or have different sample prompts per language).
- **Error Handling and User Messaging:** If the LLM fails to produce an answer (timeout or error), notify the user gracefully: e.g., "Code review could not be completed. Please check your network or API key." Avoid silent failures.
- **Performance Considerations:** Large code and complex analysis could freeze the extension host if done synchronously. Always perform LLM calls and processing asynchronously (which our example does by `await` in the command). This prevents blocking the editor UI. Also consider debouncing frequent triggers (if you ever add an auto-run, ensure it doesn't spam the API on every keystroke).

By paying attention to these UX aspects, the "Deep Code Reviewer" will feel like an intelligent assistant rather than a hindrance.



## 8. Testing, Debugging, and Deployment Best Practices

**Testing the Extension Functionality:** Develop a set of test cases – code snippets or files with known issues – to validate your extension:

- Manually run the extension on these snippets and verify it catches the intended issues (and doesn't flag false ones). For example, create a small Python file with a deliberate bug (like using `==` instead of `=` in an if in Python) and see that the extension (with the LLM) marks it.
- Automated testing: VS Code allows running integration tests for extensions ([Your First Extension | Visual Studio Code Extension API](#)). You can write a test script that opens a document, executes your command, and checks that diagnostics appear. The VS Code Extension Test Runner (Mocha-based) can be configured to run such tests (see the official **Testing Extensions** guide ([Your First Extension | Visual Studio Code Extension API](#)) ([Your First Extension | Visual Studio Code Extension API](#)) for setup). This can be part of your CI to catch regressions.
- If your extension uses a remote API, you might mock the API for tests or use a smaller local model for test environment, to avoid external calls/flaky network in tests.

**Debugging:** Use VS Code's debugger to step through your extension code. If something isn't working (say, the command doesn't show up or the API call fails), add breakpoints or `console.log` statements. Common pitfalls include forgetting to update `package.json` contributors, or not handling promise rejections from the fetch call, etc. The debug console is your friend to inspect variables (like the exact prompt being sent, or the response from the API).

**Fine-Tuned Model Debugging:** If using a local model, it's like debugging any server application. Ensure the model server is running and reachable. Check logs from that side (maybe run it with verbose logging to see incoming requests).

**Resource Usage:** If running heavy models, monitor memory/CPU. It might be wise to not load the model until first use (lazy load) and perhaps unload if not used for a long time (to free memory) – depending on your use case.

**Deployment (Packaging the Extension):** Once satisfied, you can package your extension into a `.vsix` file and/or publish it:

- Use the VS Code Extension Manager CLI (`vsce`). Install it with `npm install -g vsce`. Then run `vsce package` in your extension folder. This will create `your-extension-name-x.y.z.vsix`. This VSIX can be installed by others or used in deployment pipelines.

- To publish to the Visual Studio Marketplace, you'll need to create a publisher account (on Azure DevOps) and use `vsce publish`. This requires a Personal Access Token. Follow the official guide for publishing ([Your First Extension | Visual Studio Code Extension API](#)). Once published, your extension can be installed by name from the VS Code Extensions panel.
- **Include Documentation:** In your README (which will be shown on the Marketplace), include usage instructions, any requirements (like “needs OpenAI API key”), and screenshots if possible of the extension in action. A good README and changelog will help adoption.
- **Licensing & Compliance:** If you use an API like OpenAI, ensure your terms of use allow this. OpenAI typically allows such use but requires that users handle API keys securely. Do not hardcode your own key into the extension – instead, require the user to provide one in settings. Also, if code privacy is a concern, clearly document that code is being sent to an API (unless using a local model). For enterprises, the local model route might be preferable for confidentiality.
- **Update Strategy:** Plan how you will handle updates (especially if model improvements are made). Users will appreciate release notes describing new fine-tuning updates or features in the extension. VS Code will automatically update extensions from the marketplace.

**Performance Testing:** Before full deployment, test the extension in a “real” project environment. How does it perform on a large codebase? Perhaps integrate it into a continuous integration (CI) scenario where it runs on pull requests to see if it can handle reviewing diffs of a few hundred lines quickly. Measure latency of response (for instance, GPT-4 might take N seconds for M lines of code – ensure it's acceptable for your users).

By thoroughly testing and following VS Code's extension best practices, you'll minimize runtime errors and ensure a smooth deployment to users' editors.

## 9. Real-World Use Cases and Evaluation Metrics

**Use Case 1: AI-Augmented Peer Review.** Imagine a scenario in a team: a developer opens a pull request. Before any human reviews it, they run the “Deep Code Reviewer” on the changed files. The extension, say using GPT-4, adds comments or highlights issues like a potential null pointer dereference or a mismatched unit (e.g., mixing up meters vs feet). The developer fixes these proactively, leading to a cleaner PR. In the PR description, they might even include the AI review report for the human reviewers to see. This speeds up the review by focusing human attention on deeper design issues while trivial bugs are already caught. NVIDIA's internal experiments with AI code reviewers showed that a fine-tuned LLM could automatically assign severity to issues and provide explanations, helping developers focus on critical concerns

([Fine-Tuning Small Language Models to Optimize Code Review Accuracy | NVIDIA Technical Blog](#)) ([Fine-Tuning Small Language Models to Optimize Code Review Accuracy | NVIDIA Technical Blog](#)). Over time, the AI can reduce the workload on senior reviewers.

**Use Case 2: Learning and Onboarding Tool.** A junior developer writing code can use the extension as a tutor. Each time they write a function, they trigger the code review to get immediate feedback. For example, the extension might flag inefficient algorithms or incorrect complexity, and the LLM's explanation serves as a teaching moment ("**Issue:**  $O(n^3)$  loop detected – this will be slow on large inputs. **Suggestion:** Consider optimizing with a hash map for  $O(n)$  performance."). This is like having a mentor pair-programming, and it can accelerate the learning curve. Developers often use AI coding assistants for generation; using one for *analysis* ensures they also learn from mistakes.

**Use Case 3: Continuous Integration (CI) Guard:** Beyond VS Code, your project could be extended to a CLI that runs on a git repository (reusing the same LLM logic). This could be a bot that comments on GitHub PRs with the AI-found issues. In VS Code, you could integrate this by letting the extension compare the current file with the last committed version and only send the diff to the LLM (so it focuses on new changes). This way, the extension can be used pre-commit, and the same logic can run in CI post-commit. The **evaluation metrics** here might be: percentage of bugs caught by AI before release, reduction in QA findings, etc.

**Evaluating the Effectiveness:** To measure the success of the "Deep Code Reviewer":

- **Benchmark against known issues:** Create a benchmark suite of code with known bugs. Measure how many the LLM identifies (True Positives) versus how many it misses (False Negatives). Also, check if it flags things that are not actually problems (False Positives). These metrics correspond to **recall** and **precision**. Depending on your goals, you might aim for high recall (catch as many bugs as possible, even if a few false alarms) or a balance.
- **Severity Accuracy:** If your tool categorizes issue severity, measure if it gets it right compared to human judgement. NVIDIA's study on automated code review used **severity rating prediction accuracy** as a metric, and they improved it significantly by fine-tuning a smaller model with GPT-4's guidance ([Fine-Tuning Small Language Models to Optimize Code Review Accuracy | NVIDIA Technical Blog](#)).
- **Developer Feedback:** Nothing beats real user feedback. If deploying in your team or to users, gather qualitative input. Do developers find the suggestions useful? How often do they act on the recommendations? You could instrument the extension (with consent) to report anonymized data like "issues flagged vs issues ignored" or add a feedback command ("flag this suggestion as incorrect") to continually improve prompt or model choices.
- **Comparison with Static Analysis:** Evaluate where the AI overlaps or complements traditional linters/analysis. You may find the LLM finds more logical bugs, whereas linters

find more style issues. This can guide you to focus the AI on what linters miss. For example, an LLM might catch that a piece of code doesn't handle a certain edge-case input – something a static analyzer might not know – whereas a null reference might be catchable by both.

**Continuous Improvement:** The AI code reviewer can improve over time. As you gather examples of mistakes it misses, you can fine-tune further or adjust prompts. If it hallucinates issues (says there's a bug where there isn't), analyze those cases and refine prompts to reduce false alarms (or adjust the parsing logic to discard low-confidence output).

In summary, a “Deep Code Reviewer” VS Code extension can be a transformative tool in real development workflows. By integrating a powerful LLM with the developer's editor, you get instant code review feedback, learning opportunities, and an added layer of quality assurance. With careful engineering, fine-tuning, and user-centric design, such a tool is not only feasible but highly valuable in the modern AI-assisted programming era.

## 10. Conclusion

You have now a complete blueprint for building a VS Code extension that leverages an LLM for code review. We covered setting up the development environment ([Your First Extension | Visual Studio Code Extension API](#)) ([Your First Extension | Visual Studio Code Extension API](#)), deciding on the implementation language and architecture (JavaScript vs Python bridging) ([Programmatic Language Features | Visual Studio Code Extension API](#)), integrating with leading AI models like GPT-4 or CodeT5+ for code analysis, crafting prompts and optionally fine-tuning the model for better accuracy (as evidenced by large improvements in targeted tasks through fine-tuning ([Fine-Tuning Small Language Models to Optimize Code Review Accuracy | NVIDIA Technical Blog](#))), implementing the extension commands and UI feedback (using VS Code's APIs for editor integration ([Programmatic Language Features | Visual Studio Code Extension API](#))), and best practices for testing and deployment. We also explored how this tool can be applied in real scenarios and how to measure its impact.

By following this step-by-step guide, you can build a **real-world ready AI Code Reviewer** that augments your development process. Good luck, and happy coding – with your new AI code review partner!

### Sources:

- Visual Studio Code Official Docs – *Your First Extension* (Hello World example, debugging, etc.) ([Your First Extension | Visual Studio Code Extension API](#)) ([Your First Extension | Visual Studio Code Extension API](#))
- Visual Studio Code Official Docs – *Programmatic Language Features* (Diagnostics and CodeLens APIs) ([Programmatic Language Features | Visual Studio Code Extension API](#))

([Programmatic Language Features | Visual Studio Code Extension API](#))

- Salesforce AI Research – *CodeT5+: Open Code Large Language Models* (overview of CodeT5+ capabilities) ([CodeT5+: Open Code Large Language Models](#)) ([CodeT5+: Open Code Large Language Models](#))
- NVIDIA Technical Blog – *Fine-Tuning Small LLMs to Optimize Code Review Accuracy* (benefits of fine-tuning for code reviews, 18% accuracy boost) ([Fine-Tuning Small Language Models to Optimize Code Review Accuracy | NVIDIA Technical Blog](#)) ([Fine-Tuning Small Language Models to Optimize Code Review Accuracy | NVIDIA Technical Blog](#))
- OpenAI (McAleese et al.) – *LLM Critics Help Catch LLM Bugs* (fine-tuned GPT-4 critics catching more bugs than humans) ( ) ( )
- GitHub – “A Gentle Introduction to VS Code Extension + LLM” (tutorial using local LLM server with an extension) ([GitHub - hungson175/vscode-extension-gens: A gentle introduction to VS Extension & using LLM model](#)) ([GitHub - hungson175/vscode-extension-gens: A gentle introduction to VS Extension & using LLM model](#)).